

11.20.2018 ~ 12.07.2018.

What we'll learn.

- Tensors - main ds.
- Autograds - Calc. gradients
- Validations to check (methods)
- Transfer learning

2018 Pytorch
FB Challenge


Lesson 4

(very) Rough Note
by enyiecat.

Need

- Notebook (on github)
- PyTorch 0.4+

★ Single Layer NNs.



$$y = f\left(\sum w_i x_i + b\right)$$

$$= f(XW + b)$$

$\begin{bmatrix} x_1 & \dots & x_n \end{bmatrix} \begin{bmatrix} w_1 \\ \vdots \\ w_n \end{bmatrix}$

Tensors: generalized vectors. $(3+D)$.
 vectors \rightarrow matrix \rightarrow tensor.

★ How to do it? \Rightarrow just do it like numpy!!!
 e.g. `torch.randn((size))`, `torch.exp`.
`rand_like(tensor)`

quiz: `activation(torch.dot(features, weights) + bias)`
 or maybe `((torch.mm(x, w)).sum() + bias)`

ans: `activation((features strict * weights).sum() + bias)`
 may have to ~~transpose~~ one?
 \hookrightarrow ~~benient~~.


NOTE

★ weights.view: gives new tensor (shares data)
 weights.resize_: modifies tensor 'in-place'
 weights.reshape: usually works as view()
 but sometimes clone.

inefficient!!!

Q in this case, we can use transpose(), permute()

Stacking Them up! (Vid 4-2)



Quiz (calc L1)

$$L1 = \text{activation}((\text{features} * W1) + \text{bias}) \quad [b_1, b_2]$$

$$L2 = \text{activation}((L1 * W2) + \text{bias } 2)$$

may use torch.mm (better cuz
X unexpected result)

Vid 5

Numpy \rightarrow torch : $b = \text{torch.from_numpy}(a)$

numpy \leftarrow torch : $a = b.\text{numpy}()$

* data is shared!

\therefore support cross in-place operation.

Vid 6,7: Neural Networks in PyTorch.

MNIST example. (@torchvision)

↳ Task: identify the digits!

(torchvision.datasets.MNIST)


data → Data Loader (can set batch size, shuffle)

★ Use How? for image, label in dataloader
blah blah blah.

★ how to extract batch?

dataiter = iter(dataloader)

(64, 1, 28, 28) ← images, labels = dataiter.next() (mini batch!) → (64)



★ Then, how to?

① Flatten image ($28 \times 28 \Rightarrow 784$)

images_flat = (64, 784)

Why? we only learned dense net. (FC) \Rightarrow 1D data.

② Define output

\Rightarrow 10 classes = 10 output units.

Quiz N layer

Flatten

images.view(images.shape[0], -1)

init.

w1 = torch.randn(784, 256)

b1 = torch.randn(256)

h = activation(torch.mm(inputs, w1) + b1)

↳ REPEAT N times. (last should have (4, 10)?)

The output: 10 vals of 0 or 1

need to normalize by softmax.

\Rightarrow For Proper Probability distribution.

Quiz

softmax

torch.exp(X) / torch.sum(torch.exp(X), dim=1)

64x10

64x64.

so we add

.view(-1, 1)

64x1

sum across col.

vid 8 : Building NN. with PyTorch (with class)

★ class Network(nn.Module) ★

① `--init--`

- need super.init
- can add layer like
self.some_layer = nn.Linear(256, 10)
- can add activation like.
self.sigmoid = nn.Sigmoid()

② forward → where things are done

e.g.) $z_l = \text{self.hidden}(x)$
 $z_l = \text{self.softmax}(z_l)$

★ Functional description available.

import torch.nn.functional as F

then apply it at forward

$z_l = F.softmax(\text{self.output}(x), \text{dim}=1)$

why? activation is fixed / (no weights!)

Tip: use activation that are non-linear (e.g. ReLU)

Vid 9 - How to train in PyTorch?

Universal Function Approximators.

⇒ maps input to output. (e.g. MNIST classifier)

= deep learning approximates it.

① Loss Function = How bad is it?

(Mean Sq. Err, Cross entropy)

② Gradient Descent ⇒ ↓ Loss function.

How? Back prop. (go to Lesson 2!)

⇒ get $\frac{\partial L}{\partial w_{(s)}}$ or $\frac{\partial L}{\partial b_{(s)}}$ for grad. desc.

How to in PyTorch?

Convention) L assigned to 'criterion'

+ in `nn.CrossEntropyLoss`

input is EXPECTED to be raw scores,
NOT softmax normalized!

Tip: `nn.Sequential` (List of layers)
are best for simple net.

Code

```
model = nn.Sequential( ~~~ )  
Criterion = nn.CrossEntropyLoss()  
logits = model(images)  
loss = criterion(logits, labels)
```

Tip: `nn.LogSoftMax` is popular, convenient.
(real prob ⇒ `torch.exp(output)`)
with this, use `nn.NLLLoss`

Vid 10 Autograd.

How to do actual backprop in PyTorch?
autograd does it automatically!

To enable it (just to be sure)

`x.requires_grad_(True)`

To use it

`x.backward()` \Rightarrow calc. `x.grad`.

To disable it (*% forwardprop only spd \uparrow)

With `torch.no_grad()`:

Then, after autograd?

`loss.backward()` \Rightarrow calc `loss.grad`
by param \uparrow network

* the loss fed into optimizers.

`optimizer = optim.SGD(model.parameters(), ...)`

each step $\left[\begin{array}{l} \text{optimizer.zero_grad()} \rightarrow \text{loss accumulated!} \\ \text{calc loss -- fwd prop, loss = criterion(...)} \\ \text{loss.backward()} \\ \text{optimizer.step()} \Rightarrow \text{changes params.} \end{array} \right.$

EPOCH: one pass of training set.

Vid 12 - Fashion MNIST exercise.

\Rightarrow My soln.

ReLU, 128 \rightarrow 256 \rightarrow 128 \rightarrow 32 \rightarrow 10 (NLLLoss)

Loss after 5 epochs: 0.296.

Vid 14 - Inference and Validation

To prevent overfitting.

Inference → making predictions (forward)

Validation → for testing.
(aka test set) Not in training set, can check overfitting.

How to do inference?

- LogSoftmax → Softmax: $ps = \text{torch.exp(out)}$
- Find highest val class
 $\text{top-p, top-class} = ps.\text{topk}(1, \text{dim}=1)$
top(k) column.
 - Get whether prediction == label
 $\text{equals} = \text{top-class} == \text{labels.view}(*\text{top-class.s}$
 $64 \times 1 \rightarrow \text{NOTEBOOK HAS ERROR!}$
 - Get accuracy: div. by size of minibatch.
 $\text{acc} = \text{torch.mean}(\text{equals.type(torch.FloatTensor)})$

How to Get Rid of Overfit?

Regularization e.g. Dropout.

→ Recp) turn OFF some nodes!

use) $\text{self.dropout} = \text{nn.Dropout}(p=nn)$

$X = \text{self.dropout}(\text{F.relu}(u))$

(Output layer: X Dropout!
when turning off dropout (e.g. eval)

$\text{model.eval()} \leftrightarrow \text{model.train()}$

Vid 17 - Load / Save Models

PyTorch Network's param

⇒ Saved in **state_dict**. (in dictionary)
like `model.state_dict`.
(weight, bias)

To Save.

`torch.save(model.state_dict, 'm.pth')`

To Load

`state_dict = torch.load('m.pth')`
`model.load_state_dict(state_dict)`
↳ must be created/init in advance

REQ. SAME ARCHITECTURE.

how to solve this? **SAVE architecture (custom.pth)**

`checkpoint = { 'input_size': m, 'output_size': m, 'hidden_layers':`
list of the layers
`Each.out_features for each in model.hidden_layers,`
no. of features
`'state_dict': model.state_dict() }`

`torch.save(m, m.pth)`

#loading.

`checkpoint = torch.load(filepath)`
`model = Network(`
custom defined. `checkpoint['input_size'],`
 `" ['output_size',`
 `" ['hidden_layers'])`
`model.load_state_dict(" ['state_dict'])`

Vid 18 - Load Image Data.

e.g. Smartphone Cam.

① torchvision.datasets.ImageFolder

`dataset = datasets.ImageFolder('path', transform=...)`

path: to file folder

e.g. root / dog / xxx.png
root / cat / yyy.png

② image \rightarrow Tensor: what transform do!

transform: ensemble of processes
(e.g. crop, resize, toTensor)

`transforms = transforms.Compose([
 tms.Resize(255),
 tms.CenterCrop(224),
 tms.ToTensor(),
])`

③ Pass to DataLoader (generator)

define batch size, shuffle,

`dataloader = torch.utils.DataLoader`


`(dataset, batch-size=16, shuffle=True)`

\rightarrow generator \Rightarrow put it in for loop or
make it an iterator using `iter()`
(e.g. `next(iter(data))`)

Data Augmentation.

\rightarrow Introduce 'Randomness' in data.

Help. network 'generalize' (perf.)



how?

`train_tr = transforms.Compose([
 tms.RandomRotation(30),
 tms.RandomResizedCrop(110),
 tms.RandomHorizontalFlip(),
 tms.ToTensor(),
 tms.Normalize([0.5, 0.5, 0.5],
 [0.5, 0.5, 0.5])
])`

in `test_tr`, no rotation, no flip.

just resize, center crop.
(as this is validation)

No randomness (consistent result)

Vid 70 - Pre-trained Network

(Transfer Learning)

Ex. Dog Cat Classifier.

ImageNet-trained net. (1M+ pics)
(torchvision.models)

e.g. VGG, AlexNet, Inception

size VS performance tradeoff.

Why it works?

→ such pretrained net is
very effective feature detectors.

∴ ImageNet trained info
TRANSFERRED to our dataset.

How to?

Note: most pretrained needs 224x224 pic.

* need to match color ch. distribution

ImageNet - Mean [0.485, 0.456, 0.406]

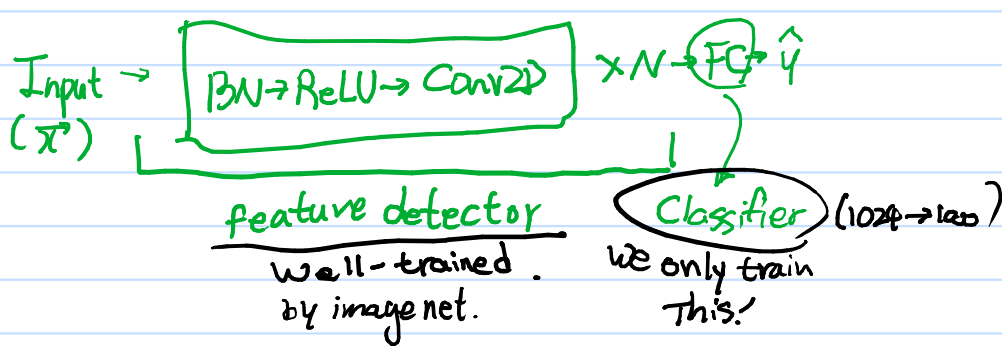
SD [0.229, 0.224, 0.225]

① Download model

from torchvision import models.

model = models.densenet121(pretrained=True)

② Look at structure!



③ Freeze feature detector part!

for param in model.parameters()
param.requires_grad = False

① x update params

② speed up the operation!

④ Build our own classifier.

classifier = nn.Sequential(

(e.g. FC(1024, 500) → ReLU
→ FC(500, 2) → ReLU → Log Softmax)

⑤ Train only classifier.

use optimizer = optim.Adam

(model.classifier.parameters(), lr)

(update only classifier)

Tip. how to run in GPU?

① model.cuda() ↔ model.cpu()

* move other tensors USED by model too!!
(e.g. images, labels)

② model.to('cuda').
vs. model.to('cpu')

check whether available by
torch.cuda.is_available()

Tip. how to format train msg?

print(f'Test loss: {blak blak : .3f} ..'
f' ~~~~~'
....)

★ Tip. what you should do? (in loop) ★

① forward prop

(don't forget optimizer.zero_grad().)

② backprop

(don't forget to store running loss)

③ Validation (for some interval)

0. model.eval()

1. fwdprop (log-ps ⇒ ps)

2. find topclass (ps.topk(1, dim=1))

3. calc. accuracy

torch.mean(equality.type(~~), item(1))
convert to Float Tensor

4. calc. loss.

(running_loss / count.)

5. model.train()

④ Print CONCISELY.

Watch those shapes

In general, you'll want to check that the tensors going through your model and other code are the correct shapes. Make use of the

```
.shape
```

method during debugging and development.

A few things to check if your network isn't training appropriately

Make sure you're clearing the gradients in the training loop with

```
optimizer.zero_grad()
```

→ accumulating gradient.

If you're doing a validation loop, be sure to set the network to evaluation mode with

```
model.eval()
```

→ dropout OFF

, then back to training mode with

```
model.train()
```

→ dropout ON

CUDA errors

Sometimes you'll see this error:

```
RuntimeError: Expected object of type torch.FloatTensor but found type torch.cuda.FloatTensor  
for argument #1 'mat1'
```

You'll notice the second type is

```
torch.cuda.FloatTensor
```

, this means it's a tensor that has been moved to the GPU. It's expecting a tensor with type

```
torch.FloatTensor
```

, no

```
.cuda
```

there, which means the tensor should be on the CPU. PyTorch can only perform operations on tensors that are on the same device, so either both CPU or both GPU. If you're trying to run your network on the GPU, check to make sure you've moved the model and all necessary tensors to the GPU with

```
.to(device)
```

where

```
device
```

is either

```
"cuda"
```

or

```
"cpu"
```